

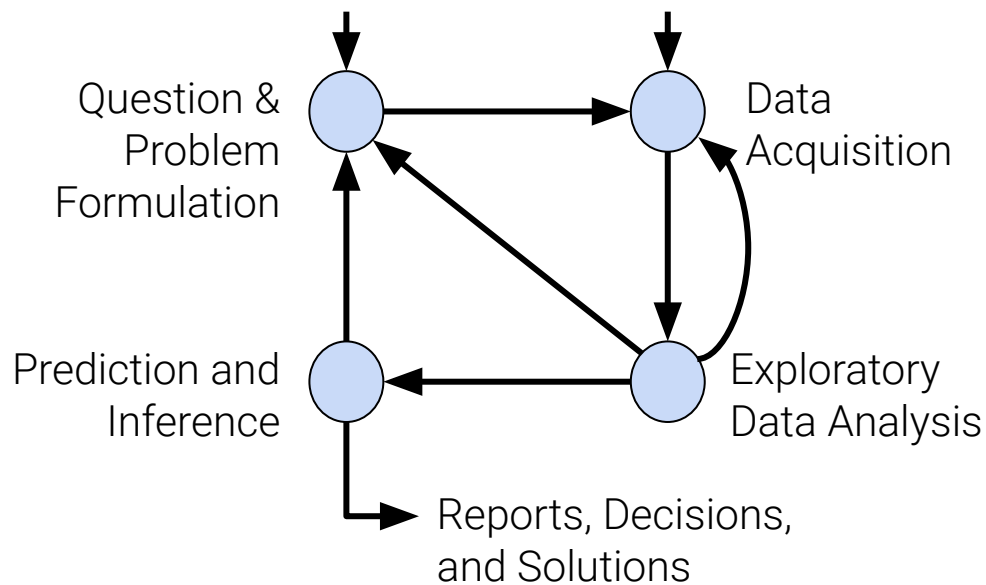
LECTURE 21

SQL I

SQL and Databases: An alternative to Pandas and CSV files.

Data 100/Data 200, Spring 2023 @ UC Berkeley

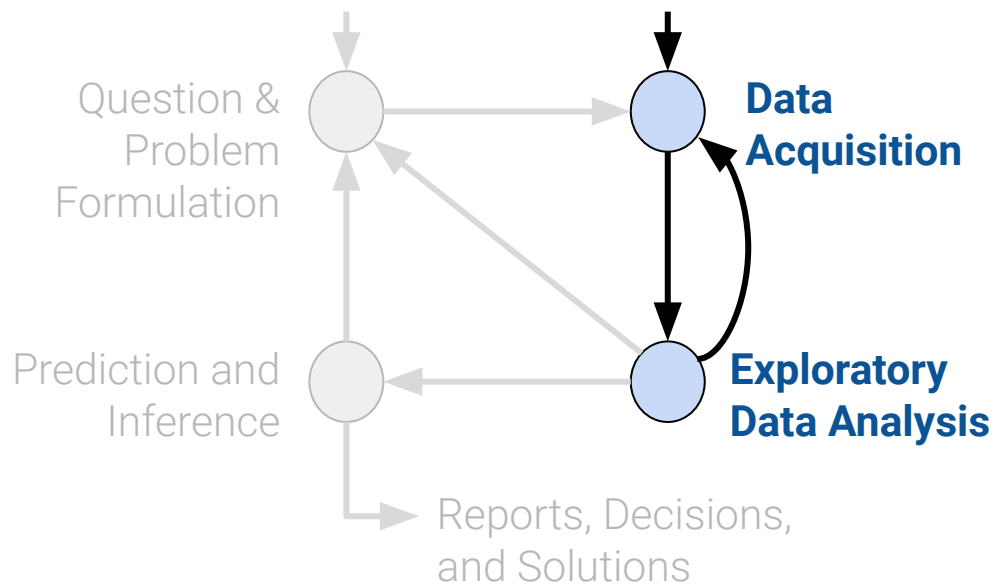
Narges Norouzi and Lisa Yan




Over the last 10 weeks, we went through the entire data science lifecycle.

In the next four weeks we are going to do it again.

- This time, with a different set of tools, ideas, and abstractions.



(today)



Today and Part of Thursday

slido



How much experience do you have with SQL?

① Start presenting to display the poll results on this slide.

Why Databases?

Lecture 21, Data 100 Spring 2023

- **Why Databases**
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- Filtering Groups with HAVING
- DISTINCT
- Python SQL



So far in Data 100, we've worked with data stored in CSV files.

Berkeley_PD_-_Calls_for_Service.csv

pd.read_csv

	CASENO	OFFENSE	EVENTDT	EVENTTM	CVLEGEND	CVDOW	InDbDate	Block_Location	BLKADDR	City	State
0	21014296	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:58	LARCENY	4	06/15/2021 12:00:00 AM	Berkeley, CA\n(37.869058, -122.270455)	NaN	Berkeley	CA
1	21014391	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:38	LARCENY	4	06/15/2021 12:00:00 AM	Berkeley, CA\n(37.869058, -122.270455)	NaN	Berkeley	CA
2	21090494	THEFT MISD. (UNDER \$950)	04/19/2021 12:00:00 AM	12:15	LARCENY	1	06/15/2021 12:00:00 AM	2100 BLOCK HASTE ST\nBerkeley, CA\n(37.864908,...	2100 BLOCK HASTE ST	Berkeley	CA
3	21090204	THEFT FELONY (OVER \$950)	02/13/2021 12:00:00 AM	17:00	LARCENY	6	06/15/2021 12:00:00 AM	2600 BLOCK WARRING ST\nBerkeley, CA\n(37.86393...)	2600 BLOCK WARRING ST	Berkeley	CA
4	21090179	BURGLARY AUTO	02/08/2021 12:00:00 AM	6:20	BURGLARY - VEHICLE	1	06/15/2021 12:00:00 AM	2700 BLOCK GARBER ST\nBerkeley, CA\n(37.86066,...	2700 BLOCK GARBER ST	Berkeley	CA

Perfectly reasonable workflow for small data that we're not actively sharing with others.



A **database** is an organized collection of data.

A **database management system (DBMS)** is a software system that **stores**, **manages**, and **facilitates access** to one or more databases.

Why use DBMSes?

- Our data might not be stored in a simple-to-read format such as a CSV (comma-separated values) file.
- Think of a CSV like an Excel sheet or a sheet in Google sheets.
- In Data 8, most of the data were given to you in CSV files, but that will not always be the case in the real world.

If our data are stored in a DBMS, we must use languages such as Structured Query Language (SQL) to query for our data.



Data Storage:

- **Reliable storage** to survive system crashes and disk failures.
- Optimize to **compute on data that does not fit in memory**.
- Special data structures to **improve performance** (see CS (W)186).

Data Management:

- Configure how data is **logically organized** and **who has access**.
- Can enforce guarantees on the data (e.g. non-negative person weight or age).
 - Can be used to **prevent data anomalies**.
 - Ensures **safe concurrent operations** on data (multiple users reading and writing simultaneously, e.g. ATM transactions).

SQL Overview

Lecture 21, Data 100 Spring 2023

- Why Databases
- **Warmup: SQL Example**
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- DISTINCT
- Python SQL

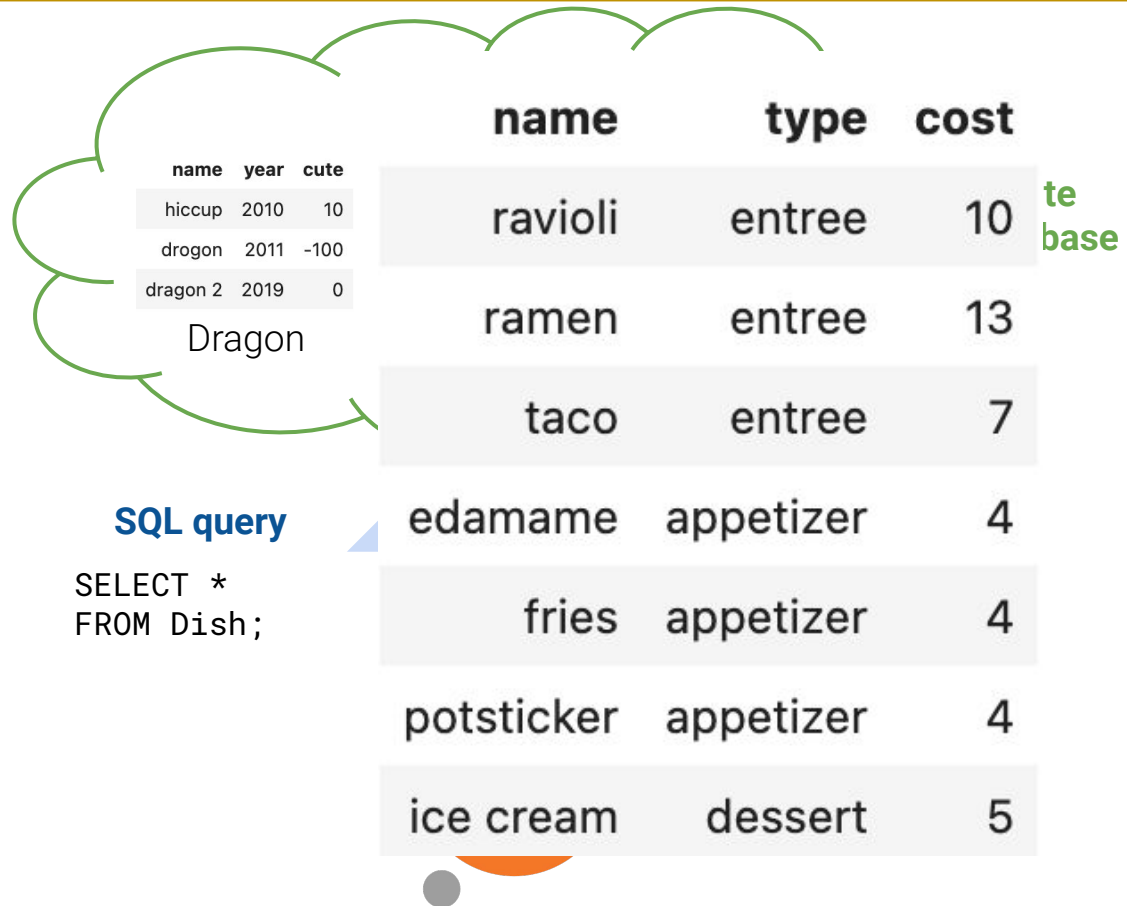


Today we'll be using a programming language called "Structured Query Language" or SQL.

- SQL is its own programming language, totally distinct from Python.
- SQL is a special purpose programming language used specifically for communicating with databases.
- We will program in SQL using Jupyter notebooks.

Let's see a quick demo of how we can use SQL to connect to a database and view a SQL Table.

Demo



Step 1: Load the SQL Module

Our first step is to load the SQL module. We do so using the `ipython cell magic` command:

```
%load_ext sql
```





Step 2: Connect to a Database

Our first step is to load the SQL module. We do so using the `ipython cell magic` command:

```
%load_ext sql
```

The second step is to connect to a database.

We use the `%%sql` header to tell Jupyter that this cell represents SQL code rather than Python code.

```
%%sql  
sqlite:///data/basic_examples.db
```

Connected: @data/18_basic_examples.db



(A note about SQLite)

Our first step is to load the SQL module. We do so using the `ipython cell magic` command:

```
%load_ext sql
```

The second step is to connect to a database.

We use the `%%sql` header to tell Jupyter that this cell represents SQL code rather than Python code.

```
%%sql  
sqlite:///data/basic_examples.db
```

Connected: @data/18_basic_examples.db

In Data 100, our database is stored in a local file. In real world practice, you'd probably connect to a remote server.

```
%%sql  
postgres://joshhug:mypassw@berkeley.edu/grades
```

There are various extensions to SQL.

We are learning the SQL commands and syntax supported by the SQLite library.

If you're curious: **SQLite** is a library that provides a relational DBMS (RDBMS). It is lightweight and offers file-based databases.



3. Run SQL Statements

Now that we're connected, let's make some queries!

For example, we might show every row in the **Dragon** table.

Thanks to the pandas magic, the resulting return data is displayed in a format almost identical to our Pandas tables (without an index).

SQL statements are terminated with semicolons. A **SQL query** is a SQL statement that returns data.

```
%%sql  
SELECT * FROM Dragon;
```



returns

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

SQL Terminology: Schema and Primary Keys

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- **SQL Tables**
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- Filtering Groups with HAVING
- DISTINCT
- Python SQL



Column or Attribute or Field

Row or
Record or
Tuple

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

← table name

SQL **tables** are also called **relations**.

SQL Style: Use *singular*, *CamelCase* names for SQL tables! For more, see [this post](#).

Column or Attribute or Field

Row or
Record or
Tuple

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

} Column Properties
ColName,
Type, Constraint

Dragon

table name

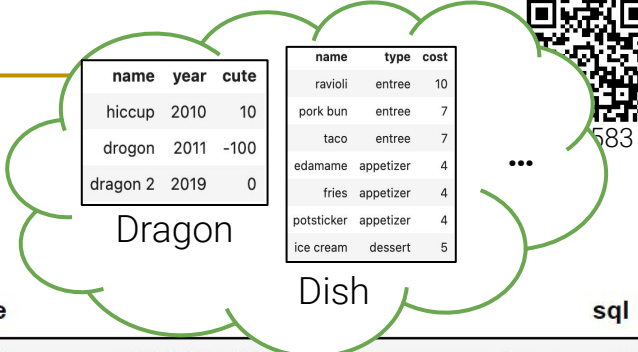
SQL **tables** are also called **relations**.

SQL Style: Use *singular, CamelCase* names for SQL tables! For more, see [this post](#).

Every column in a SQL table has three properties: **ColName**, **Type**, and zero or more **Constraints**.
(Contrast with Pandas: Series have names and types, but no constraints.)

Table Schemas

There are multiple tables in a database:



This list of tables was generated with this SQL query:

```
SELECT *
FROM sqlite_master
WHERE type='table';
```

(Many of the details here are beyond the scope of our class.)

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))
table	Scene	Scene	6	CREATE TABLE Scene (id INTEGER PRIMARY KEY AUTOINCREMENT, biome TEXT NOT NULL, city TEXT NOT NULL, visitors INTEGER CHECK (visitors >= 0), created_at DATETIME DEFAULT (DATETIME('now')))

The “**sql**” column gives the command used to create each table, and by doing so shows us each **table schema**.





Table Schemas

The table schema specifies each column schema.

Every column in a SQL table has three properties: **ColName**, **Type**, and zero or more **Constraints**.

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))
table	Scene	Scene	6	CREATE TABLE Scene (id INTEGER PRIMARY KEY AUTOINCREMENT, biome TEXT NOT NULL, city TEXT NOT NULL, visitors INTEGER CHECK (visitors >= 0), created_at DATETIME DEFAULT (DATETIME('now')))
				2 constraints
				20



Some examples of SQL **types**:

- **INT**: Integers.
- **REAL**: Real numbers.
- **TEXT**: Strings of text.
- **BLOB**: Arbitrary data, e.g. songs, video files, etc.
- **DATETIME**: A date and time.

Note: Different implementations of SQL support different types.

- SQLite: <https://www.sqlite.org/datatype3.html>
- MySQL: <https://dev.mysql.com/doc/refman/8.0/en/data-types.html>



1799583

Example Constraints

Some examples of **constraints**:

- **CHECK**: Data cannot be inserted which violates the given check constraint.
- **PRIMARY KEY**: Specifies that this key is used to uniquely identify rows in the table.
- **NOT NULL**: Null data cannot be inserted for this column.
- **DEFAULT**: Provides a value to use if user does not specify on insertion.

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))
table	Scene	Scene	6	CREATE TABLE Scene (id INTEGER PRIMARY KEY AUTOINCREMENT, biome TEXT NOT NULL, city TEXT NOT NULL, visitors INTEGER CHECK (visitors >= 0), created_at DATETIME DEFAULT (DATETIME('now')))

What is this
primary key
constraint?





Primary Keys

A **primary key** is used to uniquely identify each record in the table.

- In the Dragon table, the “**name**” of each Dragon is the primary key.
- In other words, no two dragons can have the same name!
- Primary key is used **under the hood** for all sorts of optimizations.

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Why specify primary keys?
More next time when we
discuss JOINS...



Primary Keys Can Span Columns

A **primary key** is used to uniquely identify each record in the table.

- In the Dragon table, the “**name**” of each Dragon is the primary key.
- In other words, no two dragons can have the same name!
- Primary key is used **under the hood** for all sorts of optimizations.


name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

The primary key is a **constraint**. A table can have **multiple columns** marked as primary key:

- **No duplicate tuples** allowed across the primary keys.
- Ok to have two “Serral” or two “Starcraft 2” rows.
- May only have one row with both “Serral” and “Starcraft 2”.

progamer TEXT, PK	game TEXT, PK	earnings REAL, >=0
Flash	Starcraft	580,305.13
Flash	Starcraft 2	90,152.64
Jaedong	Starcraft	418,456.82
Serral	Starcraft 2	1,143,488.55
Jaedong	Starcraft 2	224,833.53

Two video games
that have pro
esports players.



Basic SQL Queries

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- SQL Tables
- **Basic SQL Queries**
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- Filtering Groups with HAVING
- DISTINCT
- Python SQL

```
SELECT <column list>  
FROM <table>
```

;



Marks the end of a SQL statement.

Summary So Far

```
SELECT <column list>  
FROM <table>  
[WHERE <predicate>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

Goal of this section

By the end of this section, you will learn these new keywords!

But first, more SELECT

Recall our simplest query, which returns the full relation:

```
SELECT *  
FROM Dragon;
```



table name

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

We can also SELECT only a **subset of the columns**:

```
SELECT cute, year  
FROM Dragon;
```

column expression list

cute	year
10	2010
-100	2011
0	2019



WHERE: Select a rows based on conditions

To select only some rows of a table, we can use the WHERE keyword.

```
SELECT name, year
FROM Dragon
WHERE cute > 0;
```


condition

name	year
hiccup	2010

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

Mnemonic device for later:
"The Row WHERE" Dragon



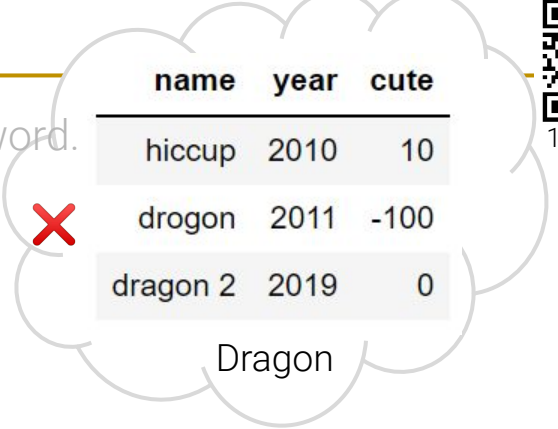
WHERE: Select a rows based on conditions

To select only some rows of a table, we can use the WHERE keyword.

```
SELECT name, year
FROM Dragon
WHERE cute > 0;
```

condition

name	year
hiccup	2010



name	year	cute
hiccup	2010	10
dragon	2011	-100
dragon 2	2019	0

Dragon




The OR, AND, and NOT let us form more complex conditions.

```
SELECT name, year
FROM Dragon
WHERE cute > 0 OR year > 2013;
```

condition

Mnemonic device for later:

“The Row WHERE” Dragon



name	year
hiccup	2010
dragon 2	2019

(fixed post-lecture. Only name, year columns returned)

ORDER BY: Sort rows

Self-explanatory

```
SELECT *  
FROM Dragon  
ORDER BY cute DESC;  
        column  
        (or ASC)
```

name	year	cute
hiccup	2010	10
dragon 2	2019	0
drogon	2011	-100

name year cute

hiccup 2010 10

drogon 2011 -100

dragon 2 2019 0

Dragon



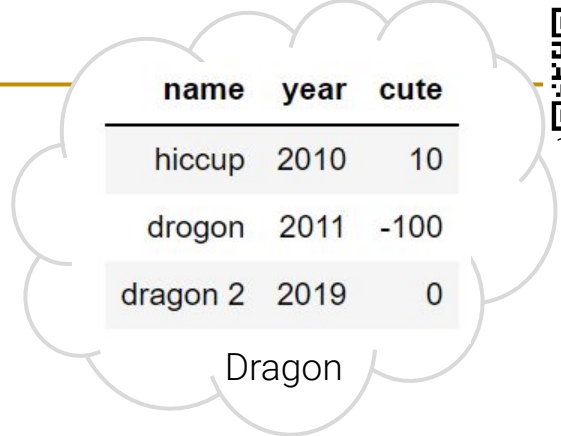
OFFSET vs. LIMIT?



1. SELECT *
FROM Dragon
LIMIT 2;

A.

name	year	cute
hiccup	2010	10
drogon	2011	-100



2. SELECT *
FROM Dragon
LIMIT 2
OFFSET 1;

B.

name	year	cute
drogon	2011	-100
dragon 2	2019	0

Matching: Which query matches each relation?

(no Slido) What do you think the **LIMIT** and **OFFSET** keywords do?



slido



Matching: Which query matches each relation?

① Start presenting to display the poll results on this slide.

OFFSET and LIMIT

The **LIMIT** keyword lets you retrieve N rows (like pandas `head()`).

```
SELECT *  
FROM Dragon  
LIMIT 2;
```

name	year	cute
hiccup	2010	10
drogon	2011	-100

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon

The **OFFSET** keyword lets you tell SQL to see later rows when limiting.

```
SELECT *  
FROM Dragon  
LIMIT 2  
OFFSET 1;
```

name	year	cute
drogon	2011	-100
dragon 2	2019	0

⚠ Unless you use **ORDER BY**, there is **no guaranteed order** of rows in the relation!



```
SELECT <column list>  
FROM <table>  
[WHERE <predicate>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

- By convention, use **all caps** for keywords in SQL statements.

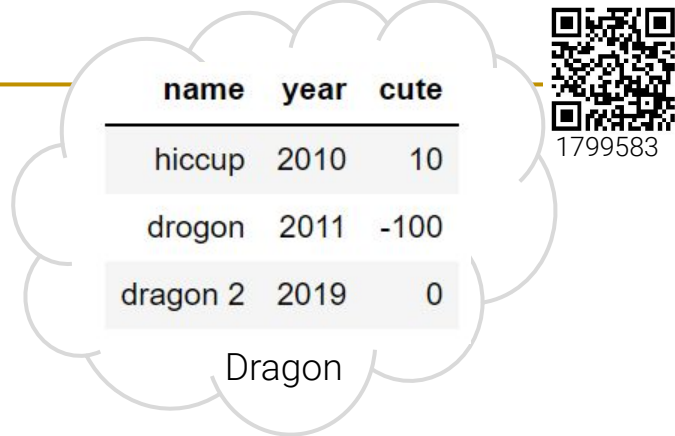
Summary So Far

The AS Keyword Aliases Columns

The AS keyword lets us rename columns during the selection process:

```
SELECT cute AS cuteness,  
       year AS birth  
FROM Dragon;
```

cuteness	birth
10	2010
-100	2011
0	2019



name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon



The AS keyword **aliases** column names.



The following two queries both retrieve the same relation:

```
SELECT cute AS cuteness,  
        year AS birth  
FROM Dragon;
```

(more readable)



cuteness	birth
10	2010
-100	2011
0	2019



```
SELECT cute AS  
cuteness, year AS  
birth FROM Dragon;
```

Use newlines and whitespace wisely in your SQL queries. It will simplify your debugging process!

```
SELECT <column expression list>  
FROM <table>  
[WHERE <predicate>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

Summary So Far

- By convention, use **all caps** for keywords in SQL statements.
- Use **newlines** to make SQL code more readable.
- **AS keyword: rename columns during selection process.**

Basic GROUP BY Operations

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- **Basic GROUP BY Operations**
- Trickier GROUP BY Operations
- Filtering Groups with HAVING
- DISTINCT
- Python SQL



```
SELECT type
FROM Dish;
```

type is not a SQL
keyword.

type

entree

entree

entree

appetizer

appetizer

appetizer

dessert

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

How do we query for a relation that **groups together** dishes of the same type?

GROUP BY

GROUP BY is similar to pandas `groupby()`.

```
SELECT type  
FROM Dragon  
GROUP BY type;
```

type

appetizer

dessert

entree

name

type

cost

ravioli

entree

10

ramen

entree

13

taco

entree

7

edamame

appetizer

4

fries

appetizer

4

potsticker

appetizer

4

ice cream

dessert

5

Dish



Aggregate Functions in Column Expression Lists

SQL has **aggregate functions**: MAX, SUM, etc.
Similar to pandas `groupby().max()`, etc.

```
SELECT type,  
       MAX(cost)  
FROM Dish  
GROUP BY type;
```

type	MAX(cost)
appetizer	4
dessert	5
entree	13

```
SELECT type,  
       SUM(cost)  
FROM Dish  
GROUP BY type;
```

type	SUM(cost)
appetizer	12
dessert	5
entree	30

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

For more aggregation functions see
https://www.sqlite.org/lang_aggfunc.html



Using Multiple Aggregation Functions

```
SELECT type,  
       SUM(cost),  
       MIN(cost),  
       MAX(name)  
FROM Dish  
GROUP BY type;
```


What do you think will happen?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



(no slido)



Using Multiple Aggregation Functions

```
SELECT type,  
       SUM(cost),  
       MIN(cost),  
       MAX(name)  
FROM Dish  
GROUP BY type;
```



type	SUM(cost)	MIN(cost)	MAX(name)
appetizer	12	4	potsticker
dessert	5	5	ice cream
entree	30	7	taco

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

No simple equivalent
in pandas!



```
SELECT <column expression list>
FROM <table>
[WHERE <predicate>]
[GROUP BY <column list>]
[ORDER BY <column list>]
[LIMIT <number of rows>]
[OFFSET <number of rows>];
```

Summary So Far

- By convention, use **all caps** for keywords in SQL statements.
- Use **newlines** to make SQL code more readable.
- **AS** keyword: rename columns during selection process.
- **Column Expressions may include aggregation functions (MAX, MIN, etc.)**

Trickier GROUP BY Operations

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- **Trickier GROUP BY Operations**
- Filtering Groups with HAVING
- DISTINCT
- Python SQL



COUNT(colname) vs COUNT(*)

```
SELECT type, COUNT(cost)
FROM Dish
GROUP BY type;
```

similar to pandas
groupby().count()

type	COUNT(cost)
appetizer	3
dessert	1
entree	3

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

```
SELECT type, COUNT(*)
FROM Dish
GROUP BY type;
```

similar to pandas
groupby().size()

type	COUNT(*)
appetizer	3
dessert	1
entree	3

COUNT(*) returns the number of rows in each group, including rows with **nulls**.

What if we GROUP BY **without** specifying enough aggregation functions in the select expression list?

```
SELECT type, cost
FROM Dish
GROUP BY type;
```

Implementation dependent:

- In some variants of SQL, this is allowed.
- In other variants, it is a syntax error.

MS SQL:

Column Dish.cost' is invalid in the select list because it is not contained in either an aggregate function or the GROUP BY clause.

This is considered bad practice. Avoid!!!

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



GROUP BY to get unique tuples?

Like in pandas, you can GROUP BY multiple columns.

```
SELECT type, cost
FROM Dish
GROUP BY type, cost;
```

type	cost
appetizer	4
dessert	5
entree	7
entree	10
entree	13

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



This works because **# cols selected = # cols to group on**, but it is a bit unwieldy. We'll learn another approach soon...

Exercise: GROUP BY Multiple Columns

- A. SELECT type, cost
FROM Dish
GROUP BY type, cost, **COUNT(*)**;
- B. SELECT type, cost, **COUNT(*)**
FROM Dish
GROUP BY type, cost;

How would we add a third column giving us the number of rows that match each type/cost tuple?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



slido



How would we add a third column giving us the number of rows that match each type/cost tuple?

① Start presenting to display the poll results on this slide.

Exercise: GROUP BY Multiple Columns

B. SELECT type, cost, COUNT(*)
FROM Dish
GROUP BY type, cost;

type	cost	COUNT(*)
appetizer	4	3
dessert	5	1
entree	7	2
entree	10	1

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



Filter Groups with HAVING

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- **Filtering Groups with HAVING**
- DISTINCT
- Python SQL

Recall earlier
mnemonic device:

“The Row WHERE” Dragon



Filtering Groups with HAVING

```
SELECT type, COUNT(*)  
FROM Dish  
GROUP BY type  
HAVING MAX(cost) < 8;
```

type	COUNT(*)
appetizer	3
dessert	1

similar to `groupby("type")`
`.filter(lambda f: max(f["cost"]) < 8)`




name	type	cost
risotto	entree	10
roast beef	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

Mnemonic device:

"The Group HAVING" Fish



WHERE vs. HAVING

```
SELECT type, COUNT(*)  
FROM Dish  
WHERE cost < 8  
GROUP BY type;
```

What will happen here?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

“The Group
HAVING” Fish



“The Row
WHERE”
Dragon



(no slido)



WHERE vs. HAVING

```
SELECT type, COUNT(*)  
FROM Dish  
WHERE cost < 8  
GROUP BY type;
```

type	COUNT(*)
appetizer	3
dessert	1
entree	1

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

“The Group
HAVING” Fish



“The Row
WHERE”
Dragon



Animation: HAVING vs. WHERE

To filter:

- Rows, use **WHERE**.
- Groups, use **HAVING**.

WHERE precedes **HAVING**.



“The Group
HAVING” Fish



“The Row
WHERE” Dragon

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish




➡ **SELECT ...**
WHERE ...
GROUP BY ...
HAVING ...

Animation: HAVING vs. WHERE

To filter:

- Rows, use **WHERE**.
- Groups, use **HAVING**.

WHERE precedes **HAVING**.



“The Group
HAVING” Fish



“The Row
WHERE” Dragon

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



1799583


SELECT ...
➡ WHERE ...
GROUP BY ...
HAVING ...

Animation: HAVING vs. WHERE

To filter:

- Rows, use **WHERE**.
- Groups, use **HAVING**.

WHERE precedes **HAVING**.



“The Group
HAVING” Fish



“The Row
WHERE” Dragon

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



SELECT ...
WHERE ...
➡ GROUP BY ...
HAVING ...

Animation: HAVING vs. WHERE

To filter:

- Rows, use **WHERE**.
- Groups, use **HAVING**.

WHERE precedes **HAVING**.



“The Group
HAVING” Fish



“The Row
WHERE” Dragon

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
spring rolls	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish



SELECT ...
WHERE ...
GROUP BY ...
➔ HAVING ...

```
SELECT <column expression list>
FROM <table>
[WHERE <predicate>]
[GROUP BY <column list>]
[HAVING <predicate>]
[ORDER BY <column list>]
[LIMIT <number of rows>]
[OFFSET <number of rows>];
```

Summary So Far

- By convention, use **all caps** for keywords in SQL statements.
- Use **newlines** to make SQL code more readable.
- **AS** keyword: rename columns during selection process.
- **WHERE: rows; HAVING: groups. WHERE precedes HAVING.**

Got to here.
Will cover on
Thursday 4/6

DISTINCT

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- Filtering Groups with HAVING
- **DISTINCT**
- Python SQL



1799583


Order of Execution

A query is **not** evaluated according to Python operator precedence.

Generally, the order of execution of clauses within a statement are:

1. **FROM**: Retrieve the Relations.
2. **WHERE**: Filter the rows.
3. **GROUP BY**: Make groups.
4. **HAVING**: Filter the groups.
5. **SELECT**: aggregate into rows, get specific columns.
6. **DISTINCT**: ???

Let's check it out!





1799583

DISTINCT: What does this do?

```
SELECT DISTINCT type
FROM Dish
WHERE cost < 11;
```


```
SELECT DISTINCT type, cost
FROM Dish
WHERE cost < 11;
```

What does the DISTINCT keyword do? Let's use the flowchart to trace each query.

type
entree
appetizer
dessert

type	cost
entree	10
entree	7
appetizer	4
dessert	5

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5





DISTINCT: What does this do?


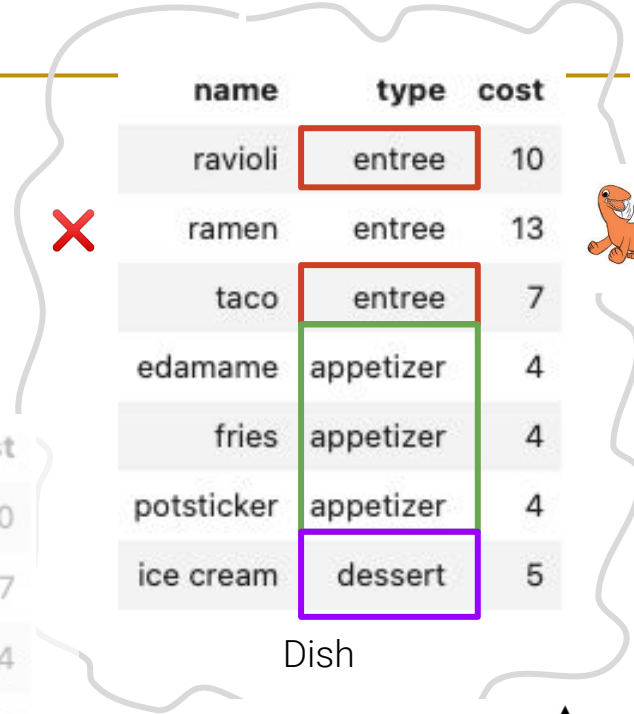
```
SELECT DISTINCT type
FROM Dish
WHERE cost < 11;
```

```
SELECT DISTINCT type, cost
FROM Dish
WHERE cost < 11;
```

type
entree
appetizer
dessert

type	cost
entree	10
entree	7
appetizer	4
dessert	5

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5





DISTINCT: What does this do?

```
SELECT DISTINCT type
FROM Dish
WHERE cost < 11;
```


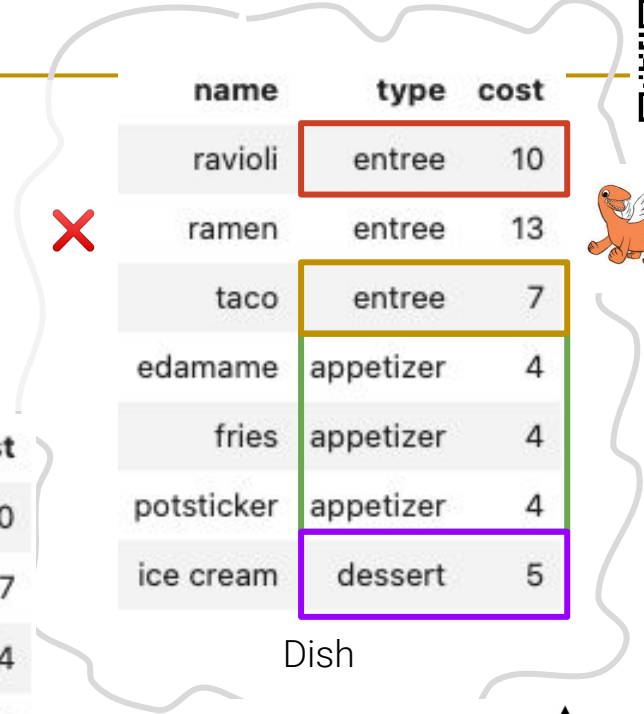
```
SELECT DISTINCT type, cost
FROM Dish
WHERE cost < 11;
```

DISTINCT creates unique tuples.
WHERE precedes **DISTINCT**.

type
entree
appetizer
dessert

type	cost
entree	10
entree	7
appetizer	4
dessert	5

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5



DISTINCT vs. GROUP BY

These queries both return unique tuples (ignoring row order) through different precedence of clauses:

```
SELECT DISTINCT
type,
cost
FROM Dish;
```


type	cost
entree	10
entree	13
entree	7
appetizer	4
dessert	5

name	type	cost
ravioli	entree	10
pork bun	entree	7
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

```
SELECT
type,
cost
FROM Dish
GROUP BY type, cost;
```

name	type	cost
ravioli	entree	10
pork bun	entree	7
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

type	cost
appetizer	4
dessert	5
entree	7
entree	10
entree	13



DISTINCT vs. GROUP BY

These queries both return unique tuples (ignoring row order) through different precedence of clauses:

```
SELECT DISTINCT
  type,
  cost
FROM Dish;
```


type	cost
entree	10
entree	13
entree	7
appetizer	4
dessert	5

Better style to use
SELECT DISTINCT for
unique values/tuples.

```
SELECT
  type,
  cost
FROM Dish
GROUP BY type, cost;
```

type	cost
appetizer	4
dessert	5
entree	7
entree	10
entree	13

I think of this as a
degenerate use of
GROUP BY, because no
aggregate functions.



☀️ DISTINCT can also be used in Column Expressions!

Common query: **GROUP BY** and **DISTINCT** together.


```
SELECT type, AVG(DISTINCT cost)
FROM DishDup
GROUP BY type;
```

type	AVG(DISTINCT cost)
appetizer	4.0
dessert	5.0
entree	8.5


← Average of the 7 and 10, which are the unique cost values for entrees.

name	type	cost
ravioli	entree	10
ramen	entree	10
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

DishDup



```
SELECT [DISTINCT] <column expression list>
FROM <table>
[WHERE <predicate>]
[GROUP BY <column list>]
[HAVING <predicate>]
[ORDER BY <column list>]
[LIMIT <number of rows>]
[OFFSET <number of rows>];
```



- By convention, use **all caps** for keywords in SQL statements.
- Use **newlines** to make SQL code more readable.
- **AS** keyword: rename columns during selection process.
- **WHERE**: rows; **HAVING**: groups. **WHERE** precedes **HAVING**.
- Column Expressions may include aggregation functions (**MAX**, **MIN**, etc.) **and DISTINCT**.

Got to here.
Will cover on
Thursday 4/6

Python-SQL

Lecture 21, Data 100 Spring 2023

- Why Databases
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- Filtering Groups with HAVING
- DISTINCT
- **Python SQL**



You can store the result of a query into a Python variable, using the << syntax:

```
%sql drag2 <<
SELECT *
FROM Dragon
LIMIT 2
OFFSET 1;
```

```
* sqlite:///data/lec18_basic_examples.db
```

```
Done.
```

```
Returning data to local variable drag2
```

This SQL result can be made into a Pandas DataFrame:

```
type(drag2)
```

```
sql.run.ResultSet
```

```
drag2.DataFrame()
```

	name	year	cute
0	dragon	2011	-100
1	dragon 2	2019	0



Both of these syntaxes work and let you use Python variables in your SQL magics:

```
query = """  
SELECT *  
FROM Dragon  
LIMIT 2  
OFFSET 1;  
"""
```

```
%%sql  
{query}
```

```
* sqlite:///data/lec18_basic_examples.db  
Done.
```

name	year	cute
drogon	2011	-100
dragon 2	2019	0

```
%%sql  
$query
```

```
* sqlite:///data/lec18_basic_examples.db  
Done.
```

name	year	cute
drogon	2011	-100
dragon 2	2019	0



Pandas knows how to talk directly to SQL engines too! You can use whichever syntax you find most convenient:

```
engine = sqlalchemy.create_engine("sqlite:///data/lec18_basic_examples.db")
connection = engine.connect()
```

```
query = """
SELECT *
FROM Dragon
LIMIT 2
OFFSET 1;
"""

pd.read_sql(query, engine)
```

	name	year	cute
0	drogon	2011	-100
1	dragon 2	2019	0


Extra SQL
practice...

Extra Slides (Sp19))

Lecture 21, Data 100 Spring 2023


- Why Databases
- Warmup: SQL Example
- SQL Tables
- Basic SQL Queries
- Basic GROUP BY Operations
- Trickier GROUP BY Operations
- DISTINCT
- **Extra Slides**

```
SELECT dept, AVG(gpa) AS avg_gpa, COUNT(*) AS size
FROM students
WHERE gender = 'F'
GROUP BY dept
HAVING COUNT(*) > 2
ORDER BY avg_gpa DESC
```



What does this compute?

```
SELECT dept, AVG(gpa) AS avg_gpa, COUNT(*) AS size
FROM students
WHERE gender = 'F'
GROUP BY dept
HAVING COUNT(*) > 2
ORDER BY avg_gpa DESC
```



What does this compute?

- The average GPA of female students and number of female students in each department where there are at least 3 female students in that department. The results are ordered by the average GPA.

```
SELECT ????
FROM tips
WHERE ????
GROUP BY ????
HAVING ????
ORDER BY ????

```

	index	total_bill	tip	sex	smoker	day	time	size
0	0	16.99	1.01	Female	No	Sun	Dinner	2
1	1	10.34	1.66	Male	No	Sun	Dinner	3
2	2	21.01	3.50	Male	No	Sun	Dinner	3
3	3	23.68	3.31	Male	No	Sun	Dinner	2
4	4	24.59	3.61	Female	No	Sun	Dinner	4

Suppose we want to compare smoker vs. non-smoker and female vs. male tips for weekend diners. Create a table ordered by percentage tip that gives the average tip for all four possibilities.

	sex	smoker	pct
0	Male	Yes	0.151530
1	Female	No	0.157180
2	Male	No	0.159930
3	Female	Yes	0.179239

```

SELECT sex, smoker, avg(tip/total_bill) as pct
  FROM tips
 WHERE day = 'Sun' OR day = 'Sat'
 GROUP BY sex, smoker
HAVING
 ORDER BY pct

```

	sex	smoker	pct
0	Male	Yes	0.151530
1	Female	No	0.157180
2	Male	No	0.159930
3	Female	Yes	0.179239

Suppose we want to compare smoker vs. non-smoker and female vs. male tips for weekend diners. Create a table ordered by percentage tip that gives the average tip for all four possibilities.

LECTURE 21

SQL I

Content credit: [Acknowledgments](#)